

AIDA roles vs Claude Code subagents (/agents)

AIDA positioning · best-effort comparison — live source:
github.com/joemooney/aida

rendered 2026-07-10

Last updated: 2026-07-09 — Claude Code’s /agents surface (the catalog of “Software Architect / Code Writer / Code Reviewer” presets, the tool-allowlist + model + prompt schema, the per-project .claude/agents/ directory) is evolving. Re-verify against the current Claude Code docs before treating capability claims here as authoritative; if the surface drifts in a way that changes the layer story, update this doc.

The TL;DR: **Claude Code subagents are a within-conversation primitive — a specialized prompt + tool allowlist + isolated context window that ends with the chat. AIDA’s roles are a cross-conversation workflow layer — full claude processes in their own git worktrees, holding leases, tied to specs in the requirement graph, with state that survives every conversation ending.** Different layers, not duplication. They compose: an AIDA role can spawn a subagent internally; AIDA is the layer above.

The worry that prompts this doc is a fair one. Claude Code’s /agents ships presets named “Software Architect,” “Code Writer,” “Code Reviewer.” AIDA ships roles named planner, implementer, reviewer. A reasonable reader sees that and asks: *am I getting the same thing twice?* The honest answer is **no** — but the rhyme is real, and saying so out loud is what makes the layer distinction stick.

What each one actually is

	Claude Code subagents (/agents)	AIDA roles + orchestrator
What it is	A specialized prompt + tool allowlist + model, with a fresh context window	A position in a lifecycle (implementer / reviewer / advisor / planner) held by a full claude process
Where it lives	Inside one Claude Code conversation; configured via .claude/agents/*.md or globally	A claude process in its own git worktree, with an aida session lease

	Claude Code subagents (/agents)	AIDA roles + orchestrator
Lifetime	Ephemeral — context window dies when the conversation ends; nothing persists	Cross-conversation — the lease, the queue entry, the spec status, and the trace comments survive every session ending
Anchoring	None — a subagent doesn't know about your requirements; it knows the prompt you gave it	Anchored to a SPEC-ID in the requirement graph (queue routing, lease tracking, status transitions, trace comments)
State	None outside the chat	Persisted in git (orphan aida-store branch, session manifests, lease files, trace comments in source)
Concurrency	Multiple subagents can run in parallel within one chat	Multiple roles run in parallel across separate worktrees, coordinated by leases + the queue
Discoverability	Visible only inside the originating conversation	Queryable via aida queue list, aida session leases, aida role list, MCP
Cost shape	Token spend inside the host conversation	Token spend across separate claude processes; coordination cost is the queue + lease infrastructure (one-time)
The shape	A tool you delegate to	A workflow position with persistent state

These are not the same kind of thing. A subagent is a *callable*; a role is a *position in a system*.

The separating test

The clearest way to feel the layer distinction is to imagine the same nominal job done two ways:

A Claude Code “Code Reviewer” subagent reviews code you hand it, in the moment; the result lives only in that conversation.

An AIDA reviewer role reviews a specific PR tied to a specific spec, writes a verdict file that drives an orchestrator merge phase, flips the spec’s status in the graph, fires the auto-bump — durable, queryable, surviving every conversation ending.

Both produce review output. Only one of them produces output that *another* session, *another* day, can find. That difference is the layer.

Why AIDA's roles can't just BE subagents

A natural follow-up: “if subagents are a built-in primitive, why doesn't AIDA just use them for its roles?” The answer is that the things AIDA's roles are made of don't exist at the subagent layer:

- **No persistence.** A subagent's context window is gone the moment the host conversation ends. You cannot build a queue out of a thing whose memory dies.
- **No graph anchoring.** A subagent doesn't know what a SPEC-ID is. Routing work to “the reviewer for STORY-86” requires a primitive that ties an actor to a graph node — subagents have no such tie.
- **No cross-session lifecycle.** AIDA's 6-phase orchestrator (commit → PR → CI → review → merge → bump) spans hours-to-days and crosses every session boundary in between. Subagents live and die inside one chat; they cannot be the substrate of a multi-day workflow.
- **No lease coordination.** Two implementers picking up the same spec from the queue is a race that the lease system prevents. Subagents share their host conversation's context — there is no “between subagents” layer to put a lease in.
- **No durable artifacts.** A reviewer subagent's verdict is a chat message. An AIDA reviewer's verdict is a file on disk that drives an orchestrator merge phase — and a structural self-merge guard blocks the implementer from merging its own work past that reviewer. Different substrates, different leverage.

AIDA's defensible niche **is** the persistent graph + the cross-session lifecycle. You cannot build it out of stateless task-runners — wrong layer. The queue, the leases, the trace comments, the auto-bump, the MCP exposure — none of these are things a within-conversation primitive can carry.

They compose

The complementary picture: **an AIDA role can spawn a subagent internally.** A Claude process held by the implementer role is itself a Claude Code session — it can invoke /agents like any other session. Concrete patterns:

- An **implementer** running a “Code Reviewer” subagent against the PR diff as a pre-/aida-pr self-check, before the AIDA reviewer role ever sees it.
- A **planner** delegating “scan this directory for affected callsites” to a subagent with a tight tool allowlist, keeping the host context clean.
- An **advisor** spawning a “Software Architect” subagent for a quick brainstorm without polluting the long-running advisory conversation.

In every case, the subagent is **a tool the role uses**, not a replacement for the role. Subagents make AIDA roles *better* — sharper focus, cleaner context windows, lower token spend per turn. AIDA is the layer above them.

The right mental model is a two-level stack: **subagents** are the within-session specialization layer; **roles** are the cross-session workflow layer. Both are useful. Neither replaces the other.

The name-rhyme caveat

It needs saying out loud: **“Software Architect / Code Writer / Code Reviewer” ≈ AIDA planner / implementer / reviewer is surface rhyme, not substance.** The naming overlap is real, and it is the single biggest source of the “am I reinventing the wheel?” worry that prompts this doc.

What the rhyme captures (genuinely): both systems carve up code work into roughly similar functional buckets. Planning is a different mental mode from writing, which is different from reviewing. Both systems noticed this and named the buckets.

What the rhyme hides (substantively): the **scope** of each bucket. A Claude Code “Code Reviewer” reviews whatever code you hand it in a chat. An AIDA reviewer reviews a specific PR for a specific spec, writes a verdict file, drives a merge phase, flips a status. Same word, different layer. The presence of three buckets in each system tells you nothing about whether they live at the same layer; the *lifecycle* of each bucket tells you everything.

If you only read the names, the answer to “*am I reinventing?*” looks like “yes.” If you read what each name produces and where it lives, the answer is “*no — different layers.*”

AIDA’s own discipline: don’t reinvent the within-session parts

A positioning doc is also a watch-item for the project it positions. The corollary of “AIDA is the cross-session layer” is **“AIDA must not reinvent the within-session layer that /agents already serves.”**

The concrete rule: if AIDA finds itself building “spawn a quick specialized in-context helper to do X within the current session,” **stop — that’s /agents.** AIDA’s niche is strictly the cross-session graph + lifecycle layer; below that line, defer to Claude Code primitives.

Worked-out cases of the discipline:

- **AIDA does NOT need** its own in-session specialization mechanism (a “spawn a focused sub-thread for this code-scan” feature). /agents already does it. AIDA shells out to it from inside a role’s claude process.
- **AIDA does NOT need** its own tool-allowlist scoping for in-conversation delegation. /agents already does it.
- **AIDA does NOT need** its own context-isolation primitive. /agents already does it.
- **AIDA DOES need** its own queue, lease, worktree, orchestrator, trace-graph, MCP layer — none of which /agents provides or attempts to.

The test, when adding any feature: “*does this give the session-level work a primitive Claude Code already provides, or does it give the cross-session workflow a primitive*”

nothing else does?” The first kind is yak shaving; the second kind is AIDA’s job.

Honest scope statement

AIDA’s value proposition vs Claude Code subagents is **not** *“better roles than /agents.”* The two systems live at different layers of the stack; comparing them as substitutes is the wrong frame. The defensible thing AIDA brings is *“a workflow layer above subagents that has graph anchoring, persistent state, queue routing, and lifecycle bookkeeping — none of which a within-conversation primitive can carry.”* That is a complementary capability, not a replacement.

If a Claude Code user has not yet hit the cross-session pain — *“yesterday’s session decided X and today’s session doesn’t know,” “two agents touched the same code and silently disagreed,” “the PR merged but the requirement is still in ‘in-progress’ because nobody flipped it”* — subagents alone may genuinely be all they need. The threshold for adding AIDA is not the number of subagents in the project; it is the first time a piece of work crosses a session boundary and needs the next session to know about it.

When that threshold lands, AIDA composes underneath: subagents stay where they are, working inside whichever role-held claude process needs them.

See also

- [vs-ultrareview.md](#) — same shape of comparison applied to Claude Code’s cloud review surface.
- [vs-ultraplan.md](#) — same shape of comparison applied to Claude Code’s cloud planning surface.
- [vs-karpathy-md.md](#) — the “is structured markdown enough?” question at the other end of the spectrum.
- [docs/aida/discipline/advisor-role.md](#) — the advisor role consults this directory when a user asks a “where does AIDA fit / vs X” question, rather than improvising.
- [OVERVIEW.md](#) “Public face: the TUI is the product” — the broader framing for why AIDA’s value is below the visible surface.